

Sophia Ju – 260852000

COMP 462 – HW3

Instructions on how to run all code is in src/README.md – need to pip install biopython

Q1.a) All results for this question are also in src/config.json file

Average lengths:

genic: 990.4515103338633

intergenic: 1051.5917566241412

Nucleotide frequencies:

"A": 0.2661070072218186,

"C": 0.24333380461567392,

"G": 0.22576010608985442,

"T": 0.26479908207265307

Start codons:

"AAA": 0.0,	"AAT": 0.0,	"AAG": 0.0,	"AAC": 0.0,	"ATA": 0.0,	"ATT": 0.0,	"ATG": 0.8876523 58240593 5,	"ATC": 0.0,
"AGA": 0.0,	"AGT": 0.0,	"AGG": 0.0,	"AGC": 0.0,	"ACA": 0.0,	"ACT": 0.0,	"ACG": 0.0,	"ACC": 0.0,
"TAA": 0.0,	"TAT": 0.0,	"TAG": 0.0,	"TAC": 0.0,	"TTA": 0.0,	"TTT": 0.0,	"TTG": 0.0376258 61155272 92,	"TTC": 0.0,
"TGA": 0.0,	"TGT": 0.0,	"TGG": 0.0,	"TGC": 0.0,	"TCA": 0.0,	"TCT": 0.0,	"TCG": 0.0,	"TCC": 0.0,
"GAA": 0.0,	"GAT": 0.0,	"GAG": 0.0,	"GAC": 0.0,	"GTA": 0.0,	"GTT": 0.0,	"GTG": 0.0747217 80604133 55,	"GTC": 0.0,
"GGA": 0.0,	"GGT": 0.0,	"GGG": 0.0,	"GGC": 0.0,	"GCA": 0.0,	"GCT": 0.0,	"GCG": 0.0,	"GCC": 0.0,
"CAA": 0.0,	"CAT": 0.0,	"CAG": 0.0,	"CAC": 0.0,	"CTA": 0.0,	"CTT": 0.0,	"CTG": 0.0,	"CTC": 0.0,
"CGA": 0.0,	"CGT": 0.0,	"CGG": 0.0,	"CGC": 0.0,	"CCA": 0.0,	"CCT": 0.0,	"CCG": 0.0,	"CCC": 0.0

Middle codons:

"AAA": 0.0359426 24736024 42,	"AAT": 0.019054 64072701 577,	"AAG": 0.0136468 72060776 092,	"AAC": 0.020269 45272154 9924,	"ATA": 0.0035508 64807396 6404,	"ATT": 0.031067 24379647 301,	"ATG": 0.023487 97852380 1765,	"ATC": 0.025304 54997910 7815,
"AGA": 0.0025344 88238264 4807,	"AGT": 0.011565 72003826 7384,	"AGG": 0.0008389 13993569 4016,	"AGC": 0.014034 06313473 1201,	"ACA": 0.0075728 12088105 3285,	"ACT": 0.013012 84667717 4603,	"ACG": 0.011301 13947106 4727,	"ACC": 0.020634 05764952 432,
"TAA": 0.0,	"TAT": 0.015757 06341383 1433,	"TAG": 0.0,	"TAC": 0.014271 21766752 8705,	"TTA": 0.0193643 93586179 86,	"TTT": 0.025743 36652959 027,	"TTG": 0.022616 79860740 277,	"TTC": 0.013746 89642154 783,
"TGA": 0.0,	"TGT": 0.005706 22845241 341,	"TGG": 0.0131402 97072351 492,	"TGC": 0.004138 10460289 5222,	"TCA": 0.0104928 78104183 439,	"TCT": 0.011183 36885273 6715,	"TCG": 0.009190 94811800 9386,	"TCC": 0.005838 51873601 4739,
"GAA": 0.0392756 94564321 31,	"GAT": 0.037388 13807879 0156,	"GAG": 0.0243397 98886503 002,	"GAC": 0.014484 17275820 4013,	"GTA": 0.0110591 45049842 784,	"GTT": 0.016592 75081511 7876,	"GTG": 0.028448 86415885 1592,	"GTC": 0.014487 39935048 6974,
"GGA": 0.0074469 74989069 918,	"GGT": 0.026971 08489325 626,	"GGG": 0.0086150 01395501 162,	"GGC": 0.025244 85802187 3068,	"GCA": 0.0193902 06324443 535,	"GCT": 0.021064 80771929 938,	"GCG": 0.030596 16132316 0963,	"GCC": 0.022297 36597138 9805,
"CAA": 0.0333468 31244383 715,	"CAT": 0.012793 43840193 3374,	"CAG": 0.0184448 14785536 476,	"CAC": 0.010667 11408746 3236,	"CTA": 0.0088440 89447591 269,	"CTT": 0.012788 59851350 8935,	"CTG": 0.029221 63301062 0328,	"CTC": 0.014580 97052669 279,
"CGA": 0.0052319 19386818 403,	"CGT": 0.020313 01171736 9874,	"CGG": 0.0028039 08693891 577,	"CGC": 0.017631 71353023 075,	"CCA": 0.0127998 91586499 293,	"CCT": 0.011254 35388296 1818,	"CCG": 0.010749 39219067 8697,	"CCC": 0.005817 54588617 5504

stop codons:

"AAA": 0.0,	"AAT": 0.0,	"AAG": 0.0,	"AAC": 0.0,	"ATA": 0.0,	"ATT": 0.0,	"ATG": 0.0,	"ATC": 0.0,
"AGA": 0.0,	"AGT": 0.0,	"AGG": 0.0,	"AGC": 0.0,	"ACA": 0.0,	"ACT": 0.0,	"ACG": 0.0,	"ACC": 0.0,
"TAA": 0.6438791 73290938,	"TAT": 0.0,	"TAG": 0.1664016 95813460 5,	"TAC": 0.0,	"TTA": 0.0,	"TTT": 0.0,	"TTG": 0.0,	"TTC": 0.0,
"TGA": 0.1897191 30895601 5,	"TGT": 0.0,	"TGG": 0.0,	"TGC": 0.0,	"TCA": 0.0,	"TCT": 0.0,	"TCG": 0.0,	"TCC": 0.0,
"GAA": 0.0,	"GAT": 0.0,	"GAG": 0.0,	"GAC": 0.0,	"GTA": 0.0,	"GTT": 0.0,	"GTG": 0.0,	"GTC": 0.0,
"GGA": 0.0,	"GGT": 0.0,	"GGG": 0.0,	"GGC": 0.0,	"GCA": 0.0,	"GCT": 0.0,	"GCG": 0.0,	"GCC": 0.0,
"CAA": 0.0,	"CAT": 0.0,	"CAG": 0.0,	"CAC": 0.0,	"CTA": 0.0,	"CTT": 0.0,	"CTG": 0.0,	"CTC": 0.0,
"CGA": 0.0,	"CGT": 0.0,	"CGG": 0.0,	"CGC": 0.0,	"CCA": 0.0,	"CCT": 0.0,	"CCG": 0.0,	"CCC": 0.0

Q1.b & c) Instructions on how to run all code is in src/README.md

For this section:

run “python **viterbi.py** -f fasta.file -c config.json -o output.file”

exact command I used is:

“python **viterbi.py** -f ../data/Vibrio_vulnificus.ASM74310v1.dna.toplevel.fa -c config.json -o vulnificus/predictions.gff3”

The GFF3 file with the gene predictions is **vulnificus/predictions.gff3**

Q1.d)

fraction of annotated genes that:

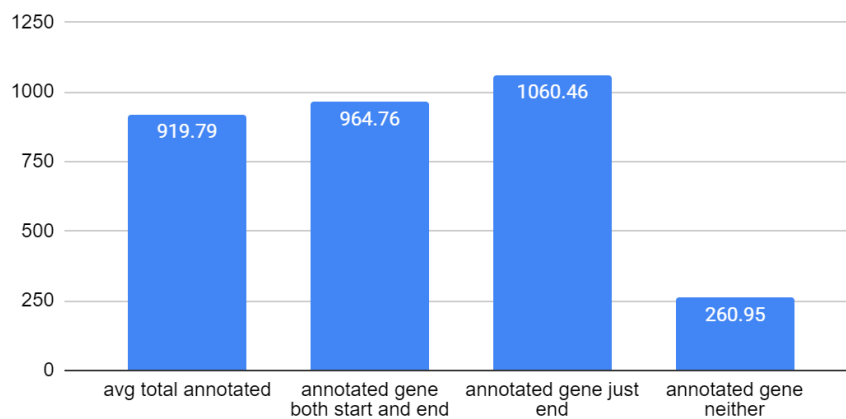
perfectly match both ends of one of predicted genes:	0.37605126151381657
match the start but not the end of a predicted gene:	0.0
match the end but not the start of a predicted gene:	0.45694833800560675
do not match either the start or end of a predicted gene:	0.16700040048057668

fraction of predicted genes that:

perfectly match both ends of one of annotated genes:	0.3626882966396292
match the start but not the end of an annotated gene:	0.0
match the end but not the start of an annotated gene:	0.44071069911162614
do not match either the start or end of an annotated gene:	0.1966010042487447

Q1.e) What properties of annotated genes are associated to an elevated risk of being partially or completely missed by your predictor?

Avg lengths of genic regions in annotated vulnificus sequences




I first calculated the average lengths of annotated genes for each category that had non-zero frequency in the previous question. Here, we can see that the annotated genes that have no match to the start or end of a predicted gene are on average much shorter in length. This is probably because we calculate transition probabilities using a much higher genic length, and so my model cannot easily find shorter genes. Additionally, from the previous question we know that the model found no annotated genes that matched the start but not the end of a predicted gene, but in many cases found just the end. This leads me to believe that when my model

predicts a gene, it will sometimes mess up finding the start of the gene especially if the gene is shorter in length, but will always be able to predict the end.

What are the properties of genes predicted by your predictor that do not match an annotated gene?

Avg lengths of genic regions in predicted vulnificus sequences



I found similar results when calculating the average lengths for predicted genes of each category. The predicted genes were about 100bp shorter on average, but followed a similar pattern as the annotated genes, suggesting that my model does okay with longer predictions but messes up often with shorter predictions. I also investigated the start and stop codon frequencies for the predicted genes that did not match the start or end of the annotated genes:

Start codons:	Emission freq	Mismatch predicted freq	Stop codons:	Emission freq	Mismatch predicted freq
ATG	0.8876	0.7075	TAA	0.6438	0.5376
TTG	0.0376	0.1785	TAG	0.1664	0.1935
GTG	0.0747	0.1140	TGA	0.1897	0.2688

From looking at the start and stop codon frequencies, mismatched predicted genes had higher frequencies of TTG or GTG as a start codon (as opposed to ATG) compared to the original frequencies from Q1.a that we used for the emission probabilities. Mismatched predicted genes also had higher frequencies of TAG and TGA as a stop codon (as opposed to TAA) compared to the emission probabilities.

Q2) For this question, I assume that since we only have distribution for up to length 1000, the maximum length of a gene region in this model is 1000bp, and that after that the model must return to the non-gene state. I also assume that since you want the distribution over the

duration of stay in the gene state to be the exact target length distribution, I should not split up the gene state by codon as we did for the HMM in Q1, and the observations are all individual nucleotides.

We could modify the given HMM by splitting the Gene state into 1000 different gene substates, one for each length from 1 to 1000. The non-gene state will only be able to transition itself and to the gene-length-1 state, and the gene-length-1 state will only be able to transition to non-gene or gene-length-2 and so on and so forth until gene-length-1000, which will only be able to transition to the non-gene state. The transition probabilities from each gene-length- n state to the non-gene state would be $\Pr[\text{length} = n] = p_n$, and the transition probabilities from a gene-length- n state to a gene-length- $(n+1)$ state would be $1 - p_n$. An exception would be for the gene-length-1000 state, which would have transmission probability = 1 to the non-gene state.

